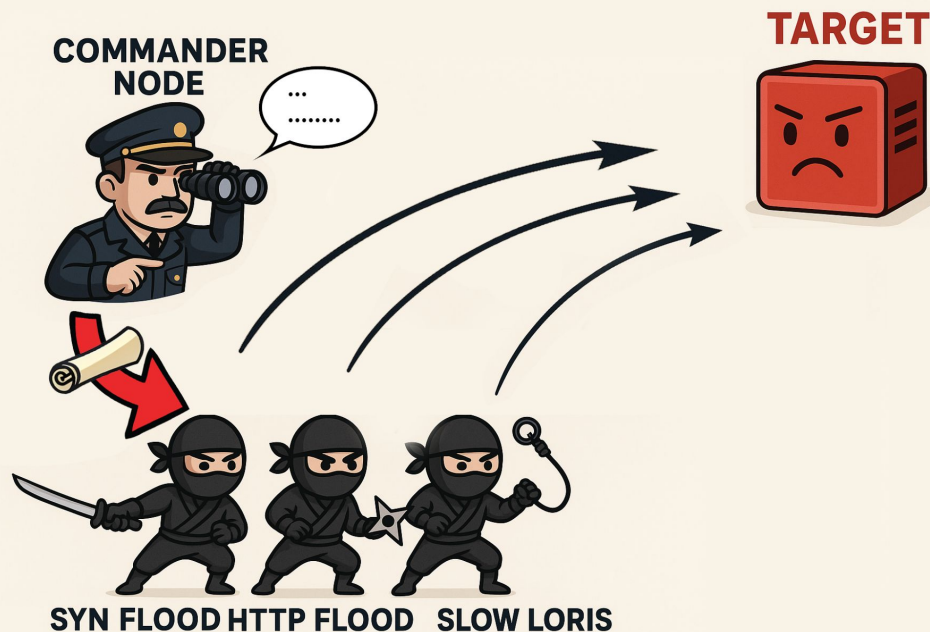


LLM-Driven DDoS Simulation: An AI-Enhanced Framework for Distributed Attack Orchestration

630 Project – Research Prototype Presentation

William Qiu



(definitely not a botnet)



Scope & Terminology

- Large Language Model(LLM)
 - Fine Tune(Fine Tuned LLM)
 - Distributed Denial of Service (DDOS)
 - Hacker: In this project/presentation, we use the term hacker to the broadest meaning(whitehat, blackhat, etc).
-
- Scope:
 - Limited Resource



Motivation & Problem Statement

Why simulate DDoS in a new way?

- **Static scripts dominate existing DDoS simulations**, relying on pre-defined behaviors with minimal ability to react or adapt.
- **Real-world attackers don't follow scripts**—they observe, learn, and adapt in real-time.
- **Existing tools are non-adaptive and unrealistic**, failing to simulate intelligent adversaries.
- **High-performance, open-source LLMs open a new possibility**: using them to emulate the intelligence, unpredictability, and strategic thinking of real attackers.
- This project explores the feasibility and implications of **AI-driven, behaviorally rich, and low-cost** simulations that mimic adversarial decision-making, potentially offering a **more accurate and challenging testbed** for cyber defense systems.



Research Gap

Existing Research Landscape

ShieldGPT: This framework utilizes LLMs for DDoS mitigation by combining traffic representation, domain knowledge, and role representation to generate mitigation strategies.

DrLLM: This approach employs LLMs in a zero-shot learning context for DDoS resistance, focusing on prompt engineering and progressive role reasoning

Similarly, **Copilot for Security**.


AI phishing attacks: voice clone, deepfake, etc.



Project Objective


- Build a research-grade LLM-based system to simulate DDoS attack patterns
- Showcase intelligent behavior coordination using a Commander/Ninja model
- Use mainstream LLMs (e.g., Mistral 7B in this case) in a resource-constrained but functional setting

System Architecture Overview





System Architecture Overview





Role of LLMs

How LLMs are embedded:

- Commander: interprets logs/status, generates attack plans
- Ninjas: execute commands, use tools (e.g., curl, slowloris), give feedback

Commander Node

- Monitoring
- LLM-empowered

```
def record_status():
    try:
        start = time.time()
        r = requests.get(TARGET, timeout=2)
        latency = round((time.time() - start) * 1000, 2)
        _window.append((r.status_code, latency))
    except Exception:
        _window.append(("ERROR", -1))

    # keep recent N records
    if len(_window) > HISTORY_LEN:
        _window.pop(0)

def summarize_observation():
    record_status()
    if not _window:
        return "No data collected."

    code_counts = {}
    total_latency = 0
```

===== LLM Query via Ollama =====

```
def query_ollama(prompt):
    try:
        response = requests.post(
            "http://localhost:11434/api/generate",
            json={
                "model": OLLAMA_MODEL,
                "prompt": prompt,
                "stream": False
            }
        )
        data = response.json()
        return data["response"].strip()
    except Exception as e:
        print(f"[COMMANDER AI] 🚨 Ollama query failed: {e}")
```

```
try:
    while True:
        obs = get_observation()
        prompt = build_prompt(obs)

        print(f"\n[COMMANDER AI] >>> Prompt to Ollama:\n{prompt}\n")

        result = query_ollama(prompt)
        decision_line = result.split("Decision:")[1].strip().upper
        valid_cmds = {"TCP", "HTTP", "SLOWLORIS"}

        if decision_line in valid_cmds:
            pub.publish(decision_line)
            print(f"[COMMANDER AI] ✅ Published: {decision_line}")
        else:
            print(f"[COMMANDER AI] ❌ Invalid command: {decision_line}")
```



Commander Node-continued

- LLM Lora Fine-tuned

```
SYSTEM_PROMPT = """You are an AI commander responsible for simulating cyberattacks based on observed server behavior.
```

```
Your goal is to:
```

1. Analyze simplified logs or status reports from a target server.
2. Decide the most appropriate type of attack from this limited set:
 - TCP
 - HTTP
 - SLOWLORIS
3. Respond only with a decision line in this format:

```
Decision: <ONE OF TCP | HTTP | SLOWLORIS>
```

```
Do not include any commentary, reasoning, or additional output. Keep it short and in log-style.
```

Ninja Node

Sub to Channel

Use DDOS-scripts

```
while True:
    msg = pubsub.get_message()
    if msg:
        print(f"[NINJA] Raw message received: {msg}")
        if msg['type'] == 'message':
            cmd = msg['data'].strip().lower()
            print(f"[NINJA] Executing received command: {cmd}")

            # swap real Dockerized Nginx server ip
            if cmd == "tcp":
                run_tcp_flood("127.0.0.1", 80, 10)
            elif cmd == "http":
                run_http_flood("http://127.0.0.1", 10)
            elif cmd == "slowloris":
                run_slowloris("127.0.0.1", 80, 30)
            else:
                print(f"[NINJA] Unknown command: '{cmd}' - ignoring.")
        time.sleep(0.1)
```

LLM-DDOS		
attack		
http_flood.py		M
slowloris.py		M
tcp_flood.py		M
infra		
> __pycache__		
monitor.py		M
pubsub.py		1, M



Demo

```
> (base) PS C:\Users\izayo\Documents\GitHub\LLM-DDOS> python -m nodes.commander
===== Commander Node Started =====
[PubSubClient] Connecting to Redis on channel 'ddos_channel'
[COMMANDER] Enter attack command (e.g., TCP, HTTP, SLOWLORIS): TCP
[PubSubClient] Publishing: TCP
[COMMANDER] Command published: TCP
[COMMANDER] Enter attack command (e.g., TCP, HTTP, SLOWLORIS):
```

- (base) PS C:\Users\izayo\Documents\GitHub\LLM-DDOS> ^C
- (base) PS C:\Users\izayo\Documents\GitHub\LLM-DDOS> python -m nodes.ninja
===== Ninja Node Started =====
[PubSubClient] Connecting to Redis on channel 'ddos_channel'
[NINJA] Waiting for commands on channel...
[PubSubClient] Subscribing to channel: ddos_channel
[PubSubClient] Subscribed. Waiting for messages...
[NINJA] Raw message received: {'type': 'subscribe', 'pattern': None, 'channel': 'ddos_channel', 'data': 1}
[NINJA] Raw message received: {'type': 'message', 'pattern': None, 'channel': 'ddos_channel', 'data': 'TCP'}
[NINJA] Executing received command: TCP

LLM-DDOS Simulation UI

A prototype control panel to configure Commanders, Ninjas, and simulate attacks on a Docker-based Nginx target.

Configure Commanders

Commander Name

1

Strategy Description

Create Commander

Output

Commander '1' created with channel '1-channel'.

List All Commanders

Commander Configs

Docker Control

Start Nginx Server

Stop Nginx Server

Docker Status

Nginx Docker container started on port 8080.

Manual Attack Trigger

Attack Command

e.g., TCP, HTTP, SLOWLORIS

Send to Ninja

Status

Configure Ninjas

Ninja Name

a

Tools (comma-separated)

Assign to Commander

1

Create Ninja

Output

Ninja 'a' created and subscribed to '1-channel'.

List All Ninjas

Ninja Configs

```
{
  "a": {
    "tools": "",
    "channel": "1-channel"
  }
}
```

Channel Mapping

Commander ↔ Ninja Channels

Show Channel Mapping



Technical Stack

Python backend


Redis (Pub/Sub system)

Open-source LLMs (Mistral 7B via Ollama)

Docker/Nginx for target simulation



gradio



redis



Future Direction

Scale Up:

- Simulate at real-world attack scale
- Deploy SuperCommander architecture for large-scale coordination

Defense Side:

- Develop adaptive defense agents based on learned attack patterns
- Explore LLM- or RL-powered agents that co-evolve with attackers

Co-Evolution Framework:

- Agents learn to attack and defend in parallel
- Periodically sync with a central brain to refine shared defense strategies